

SUNNY ZHANG

647-328-2656 | ssunny.zhang@mail.utoronto.ca | [linkedin.com/in/sunnyzhang8](https://www.linkedin.com/in/sunnyzhang8) | sunnyzhang.dev

EDUCATION

University of Toronto, St. George

Honours BSc, Double Major in Computer Science and Mathematics; GPA: 3.7/4.0

Toronto, ON

2025 – 2029

EXPERIENCE

AI Research Intern

Jun 2026 – Sep 2026

Algoverse AI Research Program

Remote

- Led a 3-person team, architected and implemented a 4-agent Llama-3.3-70B trading simulator in Python; ran 120 episodes that produced 2,400 labeled decisions.
- Built a paired-seed experiment harness for 5 pressure levels and 2 authority conditions; logged messages and capital decisions separately to compare speech with action.
- Found managers passed on capital threats in 0.3% of messages they sent versus 32.0% they received, while the trader capital gap grew from 0.55 to 0.85.

PROJECTS

LaunchPilot — Hack Canada 2x Track Winner | *Next.js, TypeScript, FastAPI, Postgres, Docker*

- Built an AI launch assistant in 36 hours: 4 agents research a market, write positioning, generate assets, queue outreach email behind a human approval gate, and paired it with a 7 day launch plan.
- Owned the agent orchestration and the inter-agent communication pipeline, threading shared project context through persistent stage threads so each agent builds on the context of the previous agents instead of restarting cold.
- Owned the auth system end to end: Auth0 JWT validation, scope-gated API routes, and the frontend session flow, plus a dev mode keeping the 49-endpoint API runnable without an Auth0 tenant.

Tiny Society | *Next.js, TypeScript, FastAPI, Supabase, Docker*

- Built a multi-agent simulator that turns any user-authored world and characters into a 7-day narrative through planning, memory retrieval, reflection, and relationship consequences.
- Architected a single-pass World Lens to extract graph facts, affordances, gathering places, and narrative register once, then ground every agent prompt and world-fit character proposal.
- Designed a deterministic 3-tier witness model to replace platform-specific action channels while preserving bounded LLM cost, legacy-save compatibility, and byte-identical batch and incremental runs.

Feedback or Annealing? | *Python, Numba, pytest, mypy*

- Benchmarked 3 aggregation rules on a 9,261-state MDP over 20 paired seeds and a matched control.
- Hardened Numba kernels with a 96-test suite run compiled, in pure Python, and bounds-checked on a cold cache, catching out-of-range reads a warm cache silently hides.
- Preregistered negative result: feedback gains at most 0.2%, and plain value iteration converges in 4.20M backups against 9.2M+ for every aggregation run.

PUBLICATIONS

1. **Sunny Zhang**, Francesco Febbo, Harshit Saini. *Language–Action Decoupling: Speech as a Weak Proxy for Action in LLM Hierarchies*. **Accepted to four NeurIPS 2026 workshops:** IAB, AIWILD, VerifyAgents, SocialAgent. sunnyzhang.dev/language-action-decoupling

TECHNICAL SKILLS

Languages: Python, TypeScript, SQL, Java, C, C++

Web: FastAPI, RestAPI, Flask, Next.js, React, PostgreSQL, SQLAlchemy, Alembic

Testing & Infrastructure: pytest, GitHub Actions, Docker, Git, Fly.io, Vercel, Supabase

ML & Data: PyTorch, NumPy, Pandas, Numba